



Global  
Address Space  
Programming Interface  
**GASPI**

# GASPI Tutorial

Christian Simmendinger

Mirko Rahn

Daniel Grünewald



# Goals

- Get an overview over GASPI
- Learn how to
  - Compile a GASPI program
  - Execute a GASPI program
- Get used to the GASPI programming model
  - one-sided communication
  - weak synchronization
  - asynchronous patterns / dataflow implementations



# Outline

- Introduction to GASPI
- GASPI API
  - Execution model
  - Memory segments
  - One-sided communication
  - Collectives
  - Passive communication



# Outline

- GASPI programming model
  - Dataflow model
  - Fault tolerance

[www.gaspi.de](http://www.gaspi.de)

[www.gpi-site.com](http://www.gpi-site.com)



Global  
Address Space  
Programming Interface  
**GASPI**  
GASPI

# Introduction to GASPI





# Motivation

- A PGAS API for SPMD execution
- Take your existing MPI code
- Rethink your communication patterns !
- Reformulate towards an asynchronous data flow model !

Bulk-synchronous Execution



Asynchronous Execution





# Key Objectives of GASPI

- **Scalability**

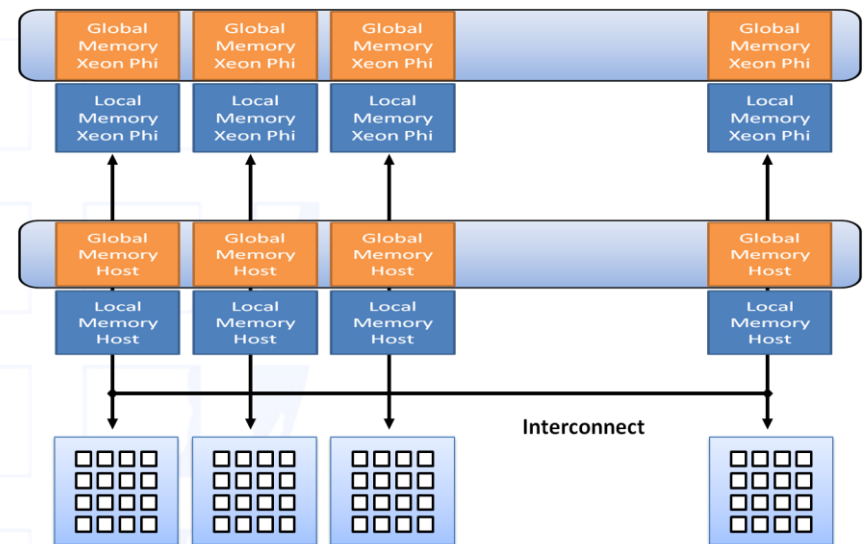
- From bulk-synchronous two sided communication patterns to asynchronous one-sided communication
- remote completion

- **Flexibility and Versatility**

- Multiple Segments,
- Configurable hardware resources
- Support for multiple memory models

- **Failure Tolerance**

- Timeouts in non-local operations
- dynamic node sets.





Global  
Address Space  
Programming Interface  
**GASPI**

# GASPI history

- **GPI**

- originally called Fraunhofer Virtual Machine (**FVM**)
- developed since 2005
- used in many of the industry projects at CC-HPC of Fraunhofer ITWM




**GPI: Winner of the „Joseph von Fraunhofer Preis 2013“**

[www.gpi-site.com](http://www.gpi-site.com)





# Scalability




## Performance

- One-sided read and writes
- **remote completion in PGAS** with notifications.
- Asynchronous execution model
  - **RDMA queues** for one-sided read and write operations, including support for arbitrarily distributed data.
- Threadsafety
  - Multithreaded communication is the default rather than the exception.
- Write, Notify, Write\_Notify
  - **relaxed synchronization** with double buffering
  - traditional (asynchronous) handshake mechanisms remain possible.
- No Buffered Communication - Zero Copy.




# Scalability




## Performance

- No polling for outstanding receives/acknowledges for send
  - **no communication overhead**, true asynchronous RDMA read/write.
- Fast synchronous collectives with time-based blocking and timeouts
  - Support for asynchronous collectives in core API.
- Passive Receives two sided semantics, no Busy-Waiting
  - Allows for distributed updates, non-time critical asynchronous collectives. Passive Active Messages, so to speak 😊.
- Global Atomics for all data in segments
  - FetchAdd
  - cmpSwap.
- Extensive profiling support.




# Flexibility and Versatility

- Segments
  - Support for **heterogeneous Memory Architectures** (NVRAM, GPGPU, Xeon Phi, Flash devices).
  - Tight coupling of Multi-Physics Solvers
  - Runtime evaluation of applications (e.g Ensembles)
- Multiple memory models
  - Symmetric Data Parallel (OpenShmem)
  - Symmetric Stack Based Memory Management
  - Master/Slave
  - Irregular.






# Flexibility



## Interoperability and Compatibility


- Compatibility with most Programming Languages.
- Interoperability with MPI.
- Compatibility with the Memory Model of OpenShmem.
- Support for all Threading Models (OpenMP/Pthreads/..)
  - similar to MPI, GASPI is orthogonal to Threads.
- GASPI is a nice match for **tile architecture** with **DMA** engines.



# Flexibility

## Flexibility

- Allows for **shrinking and growing** node set.
- User defined global reductions with **time based blocking**.
- Offset lists for RDMA read/write (write\_list, write\_list\_notify)
- **Groups** (Communicators)
- Advanced Ressource Handling, configurable setup at startup.
- Explicit connection management.



# Failure Tolerance

## Failure Tolerance.

- Timeouts in all non-local operations
- Timeouts for Read, Write, Wait, Segment Creation, Passive Communication.
- Dynamic growth and shrinking of node set.
- Fast Checkpoint/Restarts to NVRAM.
- State vectors for GASPI processes.



# The GASPI API

- 52 communication functions
  - 24 getter/setter functions
  - 108 pages
- ... but in reality:
- Init/Term
  - Segments
  - Read/Write
  - Passive Communication
  - Global Atomic Operations
  - Groups and collectives


```
GASPI_WRITE_NOTIFY ( segment_id_local  
                    , offset_local  
                    , rank  
                    , segment_id_remote  
                    , offset_remote  
                    , size  
                    , notification_id  
                    , notification_value  
                    , queue  
                    , timeout )
```

*Parameter:*

- (in) segment\_id\_local:* the local segment ID to read from
- (in) offset\_local:* the local offset in bytes to read from
- (in) rank:* the remote rank to write to
- (in) segment\_id\_remote:* the remote segment to write to
- (in) offset\_remote:* the remote offset to write to
- (in) size:* the size of the data to write
- (in) notification\_id:* the remote notification ID
- (in) notification\_value:* the value of the notification to write
- (in) queue:* the queue to use
- (in) timeout:* the timeout




## GASPI Implementation







# GASPI Implementation





Global  
Address Space  
Programming Interface

**GASPI**

www.gaspi-project.eu

# **GASPI**

## **Execution Model**



# GASPI Exection Model

- SPMD / MPMD execution model
- All procedures have prefix gaspi\_

```
gaspi_return_t  
gaspi_proc_init ( gaspi_timeout_t const timeout )
```

- All procedures have a return value
- Timeout mechanism for potentially blocking procedures



# GASPI Return Values

- Procedure return values:
  - GASPI\_SUCCESS
    - designated operation successfully completed
  - GASPI\_TIMEOUT
    - designated operation could not be finished in the given period of time
    - not necessarily an error
    - the procedure has to be invoked subsequently in order to fully complete the designated operation
  - GASPI\_ERROR
    - designated operation failed -> check error vector
- Advice: Always check return value !



# Timeout Mechanism

- Mechanism for potentially blocking procedures
  - procedure is guaranteed to return
- Timeout: `gaspi_timeout_t`
  - `GASPI_TEST (0)`
    - procedure completes local operations
    - Procedure does not wait for data from other processes
  - `GASPI_BLOCK (-1)`
    - wait indefinitely (blocking)
  - Value  $> 0$ 
    - Maximum time in msec the procedure is going to wait for data from other ranks to make progress
    - != hard execution time



# GASPI Process Management

- Initialize / Finalize
  - `gaspi_proc_init`
  - `gaspi_proc_term`
- Process identification
  - `gaspi_proc_rank`
  - `gaspi_proc_num`
- Process configuration
  - `gaspi_config_get`
  - `gaspi_config_set`



# GASPI Initialization

- `gaspi_proc_init`

```
gaspi_return_t  
gaspi_proc_init ( gaspi_timeout_t const timeout )
```

- initialization of resources

- set up of communication infrastructure if requested
- set up of default group `GASPI_GROUP_ALL`
- rank assignment

- position in machinefile  $\Leftrightarrow$  rank ID

- no default segment creation



# GASPI Finalization

- `gaspi_proc_term`

```
gaspi_return_t  
gaspi_proc_term ( gaspi_timeout_t timeout )
```

- clean up

- wait for outstanding communication to be finished
    - release resources

- no collective operation !





# GASPI Process Identification

- **gaspi\_proc\_rank**

```
gaspi_return_t  
gaspi_proc_rank ( gaspi_rank_t *rank )
```

- **gaspi\_proc\_num**

```
gaspi_return_t  
gaspi_proc_num ( gaspi_rank_t *proc_num )
```



# GASPI Process Configuration

- `gaspi_config_get`

```
gaspi_return_t  
gaspi_config_get ( gaspi_config_t *config )
```

- `gaspi_config_set`

```
gaspi_return_t  
gaspi_config_set ( gaspi_config_t const config )
```

- Retrieving and setting the configuration structure has to be done before `gaspi_proc_init`



# GASPI Process Configuration

- Configuring
  - resources
    - sizes
    - max
  - network

```
typedef struct {  
    // maximum number of groups  
    gaspi_number_t    group_max;  
  
    // maximum number of segments  
    gaspi_number_t    segment_max  
  
    // one-sided comm parameter  
    gaspi_number_t    queue_num;  
    gaspi_number_t    queue_size_max;  
    gaspi_size_t      transfer_size_max;  
  
    // notification parameter  
    gaspi_number_t    notification_num;  
  
    // passive comm parameter  
    gaspi_number_t    passive_queue_size_max;  
    gaspi_size_t      passive_transfer_size_max;  
  
    // collective comm parameter  
    gaspi_size_t      allreduce_buf_size;  
    gaspi_number_t    allreduce_elem_max;  
  
    // network selection parameter  
    gaspi_network_t    network;  
  
    // communication infrastructure build up notification  
    gaspi_number_t      build_infrastructure;  
  
    void *              user_defined;  
} gaspi_config_t;
```



# GASPI „hello world“

```
#include "success_or_die.h"
#include <GASPI.h>
#include <stdlib.h>

int main(int argc, char *argv[])
{
    SUCCESS_OR_DIE( gaspi_proc_init(GASPI_BLOCK) );

    gaspi_rank_t rank;
    gaspi_rank_t num;
    SUCCESS_OR_DIE( gaspi_proc_rank(&rank) );
    SUCCESS_OR_DIE( gaspi_proc_num(&num) );

    printf("Hello world from rank %d of %d\n",rank, num);

    SUCCESS_OR_DIE( gaspi_proc_term(GASPI_BLOCK) );
    return EXIT_SUCCESS;
}
```



# success\_or\_die.h

```
#ifndef SUCCESS_OR_DIE_H
#define SUCCESS_OR_DIE_H

#include <GASPI.h>
#include <stdlib.h>

#define SUCCESS_OR_DIE(f...) \
do \
{ \
    const gaspi_return_t r = f; \
    \
    if (r != GASPI_SUCCESS) \
    { \
        printf ("Error: '%s' [%s:%i]: %i\n", #f, __FILE__, __LINE__, r);\
        exit (EXIT_FAILURE); \
    } \
} while (0)

#endif
```



Global  
Address Space  
Programming Interface  
**GASPI**

# Hands-On





Global  
Address Space  
Programming Interface  
**GASPI**

# Memory Segments



# Segments

- Software abstraction of hardware memory hierarchy
  - NUMA
  - GPU
  - Xeon Phi
- One partition of the PGAS
- Contiguous block of virtual memory
  - no pre-defined memory model
  - memory management up to the application
- Locally / remotely accessible
  - local access by ordinary memory operations
  - remote access by GASPI communication routines






# GASPI Segments

- GASPI provides only a few relatively large segments
  - segment allocation is expensive
  - the total number of supported segments is limited by hardware constraints
- GASPI segments have an allocation policy
  - GASPI\_MEM\_UNINITIALIZED
    - memory is not initialized
  - GASPI\_MEM\_INITIALIZED
    - memory is initialized (zeroed)



# Segment Functions

- Segment creation
  - `gaspi_segment_alloc`
  - `gaspi_segment_register`
  - `gaspi_segment_create`
- Segment deletion
  - `gaspi_segment_delete`
- Segment utilities
  - `gaspi_segment_num`
  - `gaspi_segment_ptr`





# GASPI Segment Allocation

- `gaspi_segment_alloc`

`gaspi_return_t`

```
gaspi_segment_alloc ( gaspi_segment_id_t segment_id  
                      , gaspi_size_t size  
                      , gaspi_alloc_t alloc_policy )
```

- allocate and pin for RDMA

- Locally accessible

- `gaspi_segment_register`

`gaspi_return_t`

```
gaspi_segment_register ( gaspi_segment_id_t segment_id  
                        , gaspi_rank_t rank  
                        , gaspi_timeout_t timeout )
```

- segment accessible by rank



# GASPI Segment Creation

- `gaspi_segment_create`

`gaspi_return_t`

```
gaspi_segment_create ( gaspi_segment_id_t segment_id  
                      , gaspi_size_t size  
                      , gaspi_group_t group  
                      , gaspi_timeout_t timeout  
                      , gaspi_alloc_t alloc_policy )
```

- Collective short cut to

- `gaspi_segment_alloc`
- `gaspi_segment_register`

- After successful completion, the segment is locally and remotely accessible by all ranks in the group



# GASPI Segment Deletion

- `gaspi_segment_delete`

`gaspi_return_t`

```
gaspi_segment_delete ( gaspi_segment_id_t segment_id )
```

— free segment memory



# GASPI Segment Utils

- **gaspi\_segment\_num**

```
gaspi_return_t  
gaspi_segment_num ( gaspi_number_t *segment_num )
```

- **gaspi\_segment\_list**

```
gaspi_return_t  
gaspi_segment_list ( gaspi_number_t num  
                    , gaspi_segment_id_t *segment_id_list )
```

- **gaspi\_segment\_ptr**

```
gaspi_return_t  
gaspi_segment_ptr ( gaspi_segment_id_t segment_id  
                  , gaspi_pointer_t *pointer )
```



```
// includes
```

$$\{$$

```
SUCCESS OR DIE( gaspi_proc_init(GASPI_BLOCK) );
```

```
SUCCESS OR DIE( gaspi_proc_rank(&iProc));
```

```
gaspi_segment id t const segment id = 0;
```

```
gaspi_size_t      const segment_size = VLEN * sizeof (double);
```

[illegible]



## Using Segments (II)

```
gaspi_pointer_t array;  
SUCCESS_OR_DIE( gaspi_segment_ptr (segment_id, &array) );  
  
for (int j = 0; j < VLEN; ++j)  
{  
    ( (double *)array )[j]= (double)( iProc * VLEN + j );  
    printf( "rank %d elem %d: %f \n",  
           , iProc,j,( (double *)array )[j] );  
}  
  
SUCCESS_OR_DIE( gaspi_proc_term(GASPI_BLOCK) );  
return EXIT_SUCCESS;  
}
```





Global  
Address Space  
Programming Interface  
**GASPI**

# Hands-On





Global  
Address Space  
Programming Interface

**GASPI**  
1984

# One-sided Communication



# GASPI One-sided Communication

- `gaspi_write`

```
gaspi_return_t  
gaspi_write ( gaspi_segment_id_t segment_id_local  
              , gaspi_offset_t offset_local  
              , gaspi_rank_t rank  
              , gaspi_segment_id_t segment_id_remote  
              , gaspi_offset_t offset_remote  
              , gaspi_size_t size  
              , gaspi_queue_id_t queue  
              , gaspi_timeout_t timeout )
```

- Post a put request into a given queue for transferring data from a local segment into a remote segment



# GASPI One-sided Communication

- `gaspi_read`

```
gaspi_return_t  
gaspi_read ( gaspi_segment_id_t segment_id_local  
             , gaspi_offset_t offset_local  
             , gaspi_rank_t rank  
             , gaspi_segment_id_t segment_id_remote  
             , gaspi_offset_t offset_remote  
             , gaspi_size_t size  
             , gaspi_queue_id_t queue  
             , gaspi_timeout_t timeout )
```

- Post a get request into a given queue for transferring data from a remote segment into a local segment



# GASPI One-sided Communication

- `gaspi_wait`

```
gaspi_return_t  
gaspi_wait ( gaspi_queue_id_t queue  
             , gaspi_timeout_t timeout )
```

- wait on local completion of all requests in a given queue
- After successful completion, all involved local buffers are valid



# Queues (I)

- Different queues available to handle the communication requests
- Requests to be submitted to one of the supported queues
- Advantages
  - more scalability
  - channels for different types of requests
  - similar types of requests are queued and synchronized together but independently from other ones
  - separation of concerns



## Queues (II)

- Fairness of transfers posted to different queues is guaranteed
  - No queue should see its communication requests delayed indefinitely
- A queue is identified by its ID
- Synchronization of calls by the queue
- Queue order does not imply message order on the network / remote memory
- A subsequent notify call is guaranteed to be non-overtaking for all previous posts to the same queue and rank



# Weak Synchronization

- One sided-communication:
  - Entire communication managed by the local process only
  - Remote process is not involved
  - Advantage: no inherent synchronization between the local and the remote process in every communication request
- Still: At some point the remote process needs knowledge about data availability
  - Managed by weak synchronization primitives





# Weak Synchronization

- Several notifications for a given segment
  - Identified by notification ID
  - Logical association of memory location and notification



# GASPI Weak Synchronization

- `gaspi_notify`

```
gaspi_return_t  
gaspi_notify ( gaspi_segment_id_t segment_id  
               , gaspi_rank_t rank  
               , gaspi_notification_id_t notification_id  
               , gaspi_notification_t notification_value  
               , gaspi_queue_id_t queue  
               , gaspi_timeout_t timeout )
```

- posts a notification with a given value to a given queue
- remote visibility guarantees remote data visibility of all previously posted writes in the same queue, the same segment and the same process rank



# GASPI Weak Synchronization

- `gaspi_notify_waitsome`

```
gaspi_return_t  
gaspi_notify_waitsome ( gaspi_segment_id_t segment_id  
                        , gaspi_notification_id_t notific_begin  
                        , gaspi_number_t notification_num  
                        , gaspi_notification_id_t *first_id  
                        , gaspi_timeout_t timeout )
```

- monitors a contiguous subset of notification id's for a given segment
- returns successfull if at least one of the monitored id's is remotely updated to a value unequal zero



# GASPI Weak Synchronization

- `gaspi_notify_reset`


```
gaspi_return_t  
gaspi_notify_reset ( gaspi_segment_id_t segment_id  
                    , gaspi_notification_id_t notification_id  
                    , gaspi_notification_t *old_notification_val)
```

- Atomically resets a given notification id and yields the old value



# Communication example

- init local buffer
- write to remote buffer
- wait for data availability
- print





# onesided.c (I)

```
// includes

int main(int argc, char *argv[])
{
    static const int VLEN = 1 << 2;
    SUCCESS_OR_DIE( gaspi_proc_init(GASPI_BLOCK) );
    gaspi_rank_t iProc, nProc;
    SUCCESS_OR_DIE( gaspi_proc_rank(&iProc));
    SUCCESS_OR_DIE( gaspi_proc_num(&nProc));
    gaspi_segment_id_t const segment_id = 0;
    gaspi_size_t          const segment_size = 2 * VLEN * sizeof (double);
    SUCCESS_OR_DIE ( gaspi_segment_create ( segment_id, segment_size
                                           , GASPI_GROUP_ALL, GASPI_BLOCK
                                           , GASPI_MEM_UNINITIALIZED ) );

    gaspi_pointer_t array;
    SUCCESS_OR_DIE ( gaspi_segment_ptr (segment_id, &array) );
    double * src_array = (double *) (array);
    double * rcv_array = src_array + VLEN;

    for (int j = 0; j < VLEN; ++j) {
        src_array[j]= (double)( iProc * VLEN + j ); }
}
```



## onesided.c (II)

```
gaspi_notification_id_t data_available = 0;
gaspi_queue_id_t queue_id = 0;
gaspi_offset_t loc_off = 0;
gaspi_offset_t rem_off = VLEN * sizeof (double);

wait_for_queue_entries_for_write_notify ( &queue_id );
SUCCESS_OR_DIE ( gaspi_write_notify ( segment_id, loc_off
                                     , RIGHT (iProc, nProc)
                                     , segment_id, rem_off
                                     , VLEN * sizeof (double)
                                     , data_available, 1 + iProc, queue_id
                                     , GASPI_BLOCK ) );

wait_or_die (segment_id, data_available, 1 + LEFT (iProc, nProc) );

for (int j = 0; j < VLEN; ++j)
{
printf("rank %d rcv elem %d: %f \n", iProc,j,rcv_array[j] );      }
wait_for_flush_queues();
SUCCESS_OR_DIE( gaspi_proc_term(GASPI_BLOCK) );
return EXIT_SUCCESS;
}
```



# waitsome.c

```
include "waitsome.h",
#include "assert.h",
#include "success_or_die.h",

void wait_or_die ( gaspi_segment_id_t segment_id
                  , gaspi_notification_id_t notification_id
                  , gaspi_notification_t expected )
{
    gaspi_notification_id_t id;
    SUCCESS_OR_DIE (gaspi_notify_waitsome (segment_id, notification_id,
                                           1, &id, GASPI_BLOCK) );

    ASSERT (id == notification_id);

    gaspi_notification_t value;
    SUCCESS_OR_DIE (gaspi_notify_reset (segment_id, id, &value));
    ASSERT (value == expected);
}
```





# Hands-On






## Extended One-sided Calls

- `gaspi_write_notify`
  - `gaspi_write` + subsequent `gaspi_notify`
- `gaspi_write_list`
  - several subsequent `gaspi_writes` to the same rank
- `gaspi_write_list_notify`
  - `gaspi_write_list` + subsequent `gaspi_notify`
- `gaspi_read_list`
  - several subsequent `gaspi_reads`





# Matrix Transpose



**Matrix Transpose => Global Transpose + Local Transpose  
=> MPI\_Alltoall + Local Transpose**



# MPI Matrix Transpose



```
// pseudocode
#pragma omp parallel
{
    #pragma omp master
        MPI_Alltoall()
    #pragma omp barrier
    for_all_threadprivate_tiles
        do_local_transpose(tile);
}
```



# MPI - Alltoall

MPI Alltoall




Partial Results of Alltoall



Local Transpose

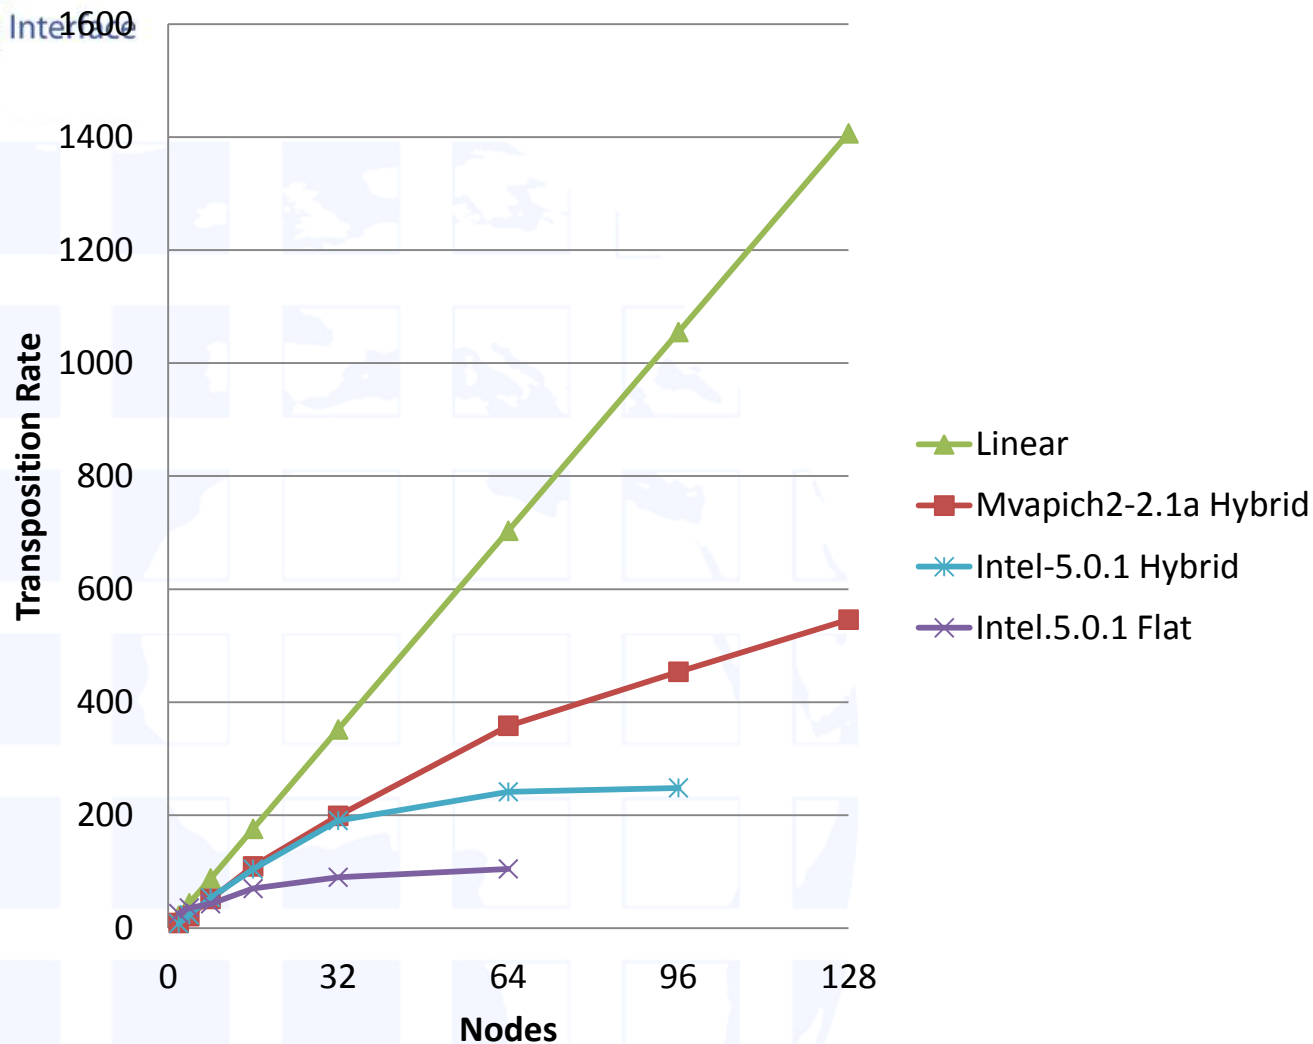


Barrier



Time





Infiniband FDR Fat Tree. Flat MPI vs. Hybrid OpenMP/MPI



Global  
Address Space  
Programming Interface

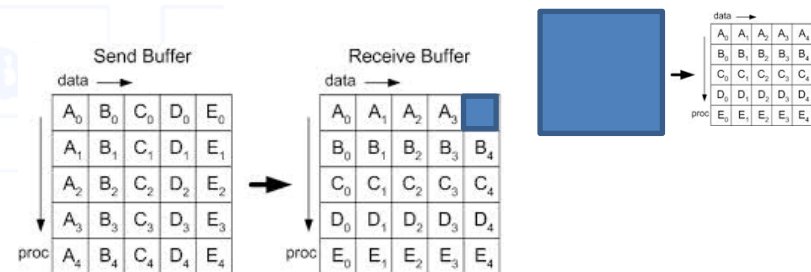
**GASPI**

# Hands-On





# GASPI Matrix Transpose




```
// pseudocode - write_notify
#pragma omp parallel
{
  #pragma omp master
  for_all_other_ranks
    gaspi_write_notify(tile)
  while (!complete)
  {
    test_or_die(thread_local tile) // test for notifications for
                                   // thread local tiles
    do_local_transpose(tile)
  }
}
```





# GASPI - Notification





# Hands-On





# MPI - GATS/PSCW

MPI Post/Start



MPI Put




MPI Complete/Wait



Local Transpose

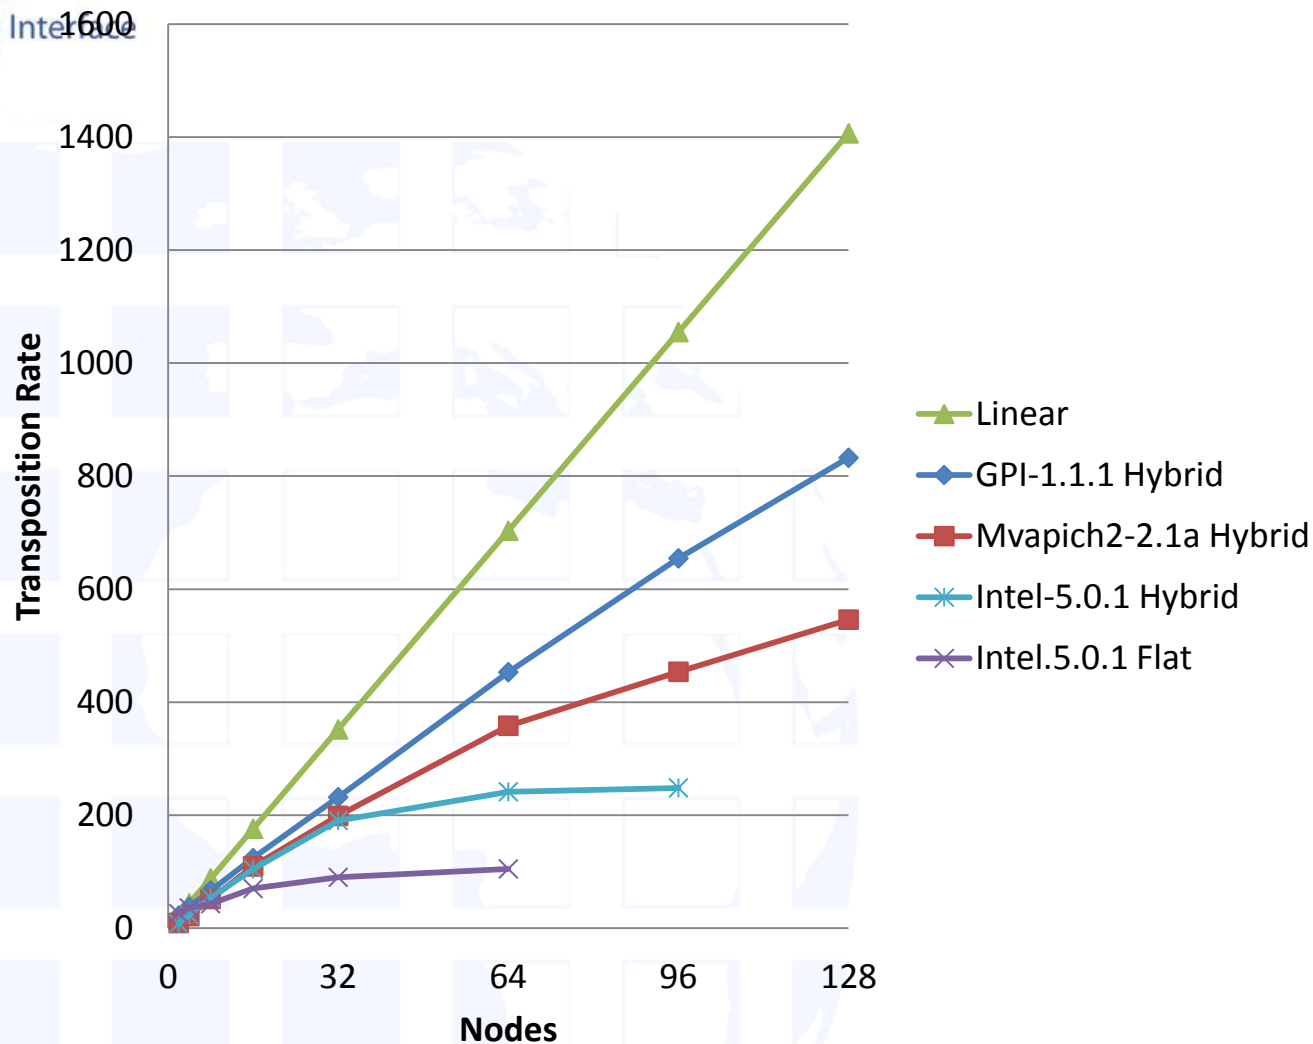


Barrier



Time







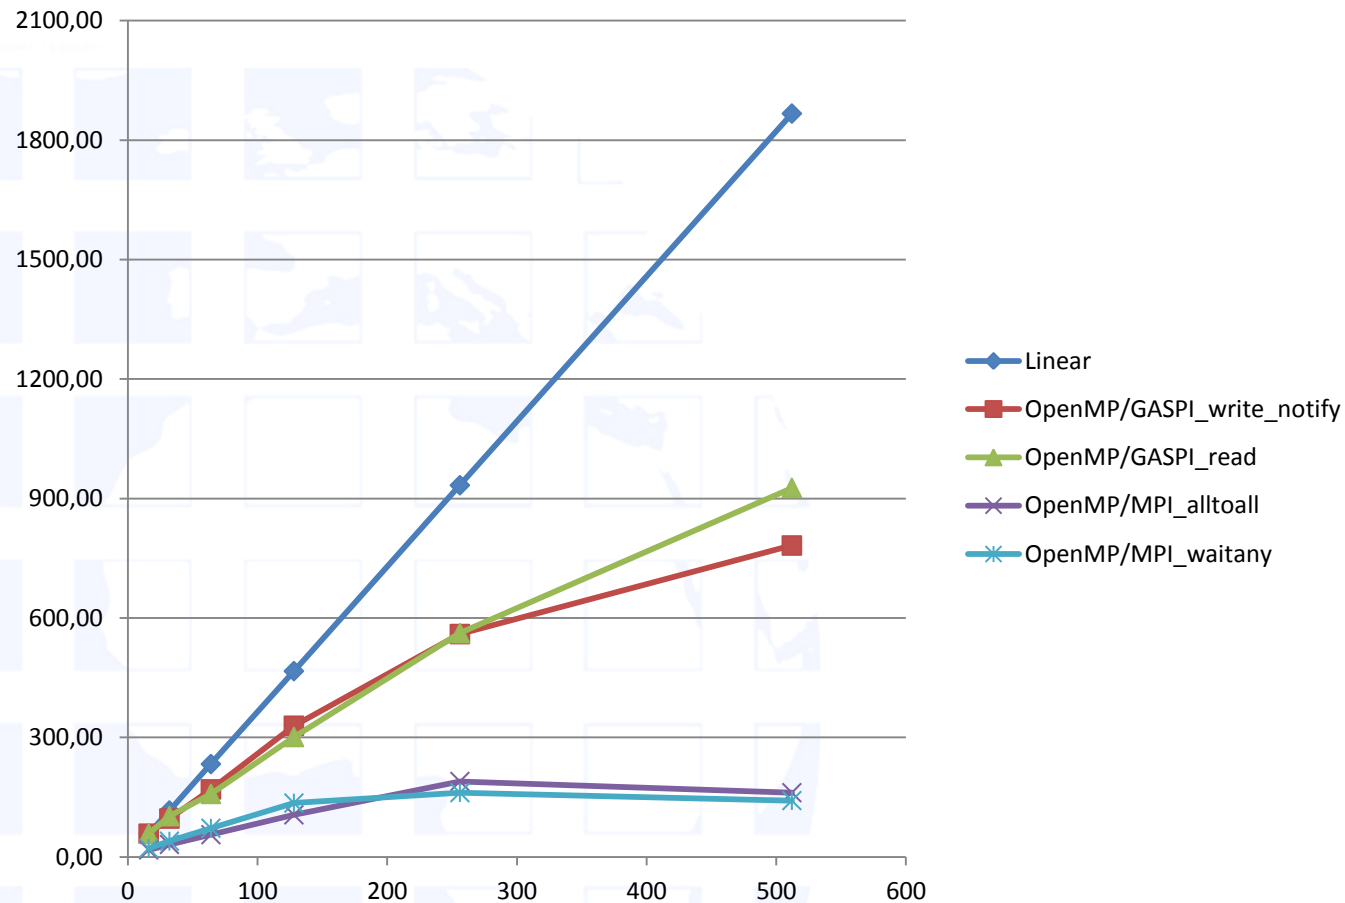
<https://github.com/PGAS-community-benchmarks>



# GASPI Matrix Transpose



```
// pseudocode - read
#pragma omp parallel
{
  #pragma omp master
  for_all_other_ranks
    gaspi_read(tile) // flip queue, wait and read !
  while (!complete)
  {
    test_or_die(thread_local tile) // test for notifications for
                                   // thread local tiles
    do_local_transpose(tile)
  }
}
```



Infiniband 7D Enhanced hypercube. GASPI read vs notified write.



Global  
Address Space  
Programming Interface

**GASPI**

# Hands-On





# Task (Graph) Models

Bottom up: Complement local task dependencies  
with remote data dependencies.

## Task (Graph) Models

### Targets

- Node local execution on (heterogeneous) manycore architectures.
- Scalability issues in Fork-Join models
- Vertically fragmented memory, separation of access and execution, handling of data marshalling, tiling, etc.
- Inherent node local load imbalance

## GASPI

### Targets:

- Latency issues, overlap of communication and computation.
- Asynchronous fine-grain dataflow model
- Fault tolerance, system noise, jitter.

Top Down: Reformulate towards asynchronous dataflow model.  
Overlap communication and computation.





Global  
Address Space  
Programming Interface

**GASPI**

# Dataflow model

Hands On Session

The MPI/GASPI Ring Exchange


Ghost Cell Exchange with Double  
Buffering



# The MPI Ring Exchange

## MPI – MPI\_Issend/MPI\_Recv

- NITER iterations of Ring Exchange with „nProc“ cores
- Shift upper half of vector to the right
- Shift lower half of vector to the left



**Example: 4 Sockets/16 cores – each core holds a vector of length  $2 \times \text{VLEN}$**



# The MPI Ring Exchange

## MPI – left\_right\_double\_buffer.c

```
for (int i = 0; i < nProc; ++i) {  
    MPI_Request send_req[2], recv_req[2];  
    const int left_halo    = 0; slice_id    = 1; right_halo    = 2;  
    MPI_Irecv ( &array_ELEM_right (buffer_id, left_halo, 0), VLEN,  
                MPI_DOUBLE, left, i, MPI_COMM_WORLD, &send_req[0]);  
    MPI_Irecv ( &array_ELEM_left (buffer_id, right_halo, 0), VLEN,  
                MPI_DOUBLE, right, i, MPI_COMM_WORLD, &send_req[1]);  
    MPI_Isend ( &array_ELEM_right (buffer_id, slice_id, 0), VLEN,  
                MPI_DOUBLE, right, i, MPI_COMM_WORLD, &recv_req[0]);  
    MPI_Isend ( &array_ELEM_left (buffer_id, slice_id, 0), VLEN,  
                MPI_DOUBLE, left, i, MPI_COMM_WORLD, &recv_req[1]);  
    MPI_Waitall (2, recv_req, MPI_STATUSES_IGNORE);  
    data_compute (NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);  
    MPI_Waitall (2, send_req, MPI_STATUSES_IGNORE);  
    buffer_id = 1 - buffer_id;  
}
```



# The MPI Ring Exchange


## MPI – left\_right\_double\_buffer\_req\_free.c

```
for (int i = 0; i < nProc; ++i) {  
    MPI_Request send_req[2], recv_req[2];  
    const int left_halo    = 0; slice_id    = 1; right_halo    = 2;  
    MPI_Irecv ( &array_ELEM_right (buffer_id, left_halo, 0), VLEN,  
                MPI_DOUBLE, left, i, MPI_COMM_WORLD, &send_req[0]);  
    MPI_Irecv ( &array_ELEM_left (buffer_id, right_halo, 0), VLEN,  
                MPI_DOUBLE, right, i, MPI_COMM_WORLD, &send_req[1]);  
    MPI_Isend ( &array_ELEM_right (buffer_id, slice_id, 0), VLEN,  
                MPI_DOUBLE, right, i, MPI_COMM_WORLD, &recv_req[0]);  
    MPI_Isend ( &array_ELEM_left (buffer_id, slice_id, 0), VLEN,  
                MPI_DOUBLE, left, i, MPI_COMM_WORLD, &recv_req[1]);  
    MPI_Request_free(&send_req[0]);  
    MPI_Request_free(&send_req[1]);  
    MPI_Waitall (2, recv_req, MPI_STATUSES_IGNORE);  
    data_compute (NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);  
    buffer_id = 1 - buffer_id;  
}
```



# The MPI Ring Exchange

- Bi-directional halo exchange –  
**implicit synchronization**





Global  
Address Space  
Programming Interface

**GASPI**

# Hands-On






# The MPI Ring Exchange

## MPI – HYBRID MPI/OpenMP

- Shift upper half of the vector to the right
- Shift lower half of the vector to the left



Example: 4 Sockets/16 cores – each core holds a vector of length  $2 \times \text{VLEN}$



# The MPI Ring Exchange

- MPI – left\_right\_double\_buffer\_funneled.c

```
for ( int i = 0; i < nProc * NTHREADS; ++i ) {  
    const int left_halo = 0, slice_id = tid + 1, right_halo = NTHREADS+1;  
    if (tid == 0) {  
        MPI_Request send_req[2], recv_req[2];  
        MPI_Irecv ( &array_ELEM_right (buffer_id, left_halo, 0), VLEN,  
                    MPI_DOUBLE, left, i, MPI_COMM_WORLD, &recv_req[0]);  
        MPI_Irecv ( &array_ELEM_left (buffer_id, right_halo, 0), VLEN,  
                    MPI_DOUBLE, right, i, MPI_COMM_WORLD, &recv_req[1]);  
        MPI_Isend ( &array_ELEM_right (buffer_id, slice_id, 0), VLEN,  
                    MPI_DOUBLE, right, i, MPI_COMM_WORLD, &send_req[0]);  
        MPI_Isend ( &array_ELEM_left (buffer_id, slice_id, 0), VLEN,  
                    MPI_DOUBLE, left, i, MPI_COMM_WORLD, &send_req[1]);  
        MPI_Request_free(&send_req[0]);  
        MPI_Request_free(&send_req[1]);  
        MPI_Waitall (2, recv_req, MPI_STATUSES_IGNORE);  
    }  
    #pragma omp barrier  
    data_compute (NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);  
    #pragma omp barrier  
    buffer_id = 1 - buffer_id; }
```






# The MPI Ring Exchange

MPI – left\_right\_double\_buffer\_funneled.c

- Fork-join model





# The MPI Ring Exchange

- MPI – `left_right_double_buffer_multiple.c`

```
if (tid == 0) {
    MPI_Request request;
    MPI_Isend ( &array_ELEM_left (buffer_id, slice_id, 0), VLEN,
               MPI_DOUBLE, left, i, MPI_COMM_WORLD, &request);
    MPI_Request_free(&request);
    MPI_Recv ( &array_ELEM_right (buffer_id, left_halo, 0), VLEN,
              MPI_DOUBLE, left, i, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    data_compute (NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
} else if (tid < NTHREADS - 1){
    data_compute (NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
} else {
    MPI_Request request;
    MPI_Isend ( &array_ELEM_right (buffer_id, slice_id, 0), VLEN,
               MPI_DOUBLE, right, i, MPI_COMM_WORLD, &request);
    MPI_Request_free(&request);
    MPI_Recv ( &array_ELEM_left (buffer_id, right_halo, 0), VLEN,
              MPI_DOUBLE, right, i, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    data_compute (NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
}
#pragma omp barrier
buffer_id = 1 - buffer_id;
```



Global  
Address Space  
Programming Interface  
**GASPI**

# Hands-On






# The GASPI Ring Exchange

## GASPI – HYBRID GASPI/OpenMP

- Shift upper half of the vector to the right
- Shift lower half of the vector to the left



Example: 4 Sockets/16 cores – each core holds a vector of length  $2 \times \text{VLEN}$



# The GASPI Ring Exchange

- GASPI – left\_right\_double\_buffer\_funneled.c

```
if (tid == 0) {
    wait_for_queue_max_half (&queue_id);
    SUCCESS_OR_DIE ( gaspi_write_notify
        ( segment_id,array_OFFSET_left(buffer_id, slice_id, 0), left,
          segment_id,array_OFFSET_left(buffer_id,right_halo,0),VLEN* sizeof(double),
          right_data_available[buffer_id], 1 + i, queue_id, GASPI_BLOCK));
    wait_for_queue_max_half (&queue_id);
    SUCCESS_OR_DIE ( gaspi_write_notify
        ( segment_id, array_OFFSET_right (buffer_id, slice_id, 0), right,
          segment_id,array_OFFSET_right(buffer_id,left_halo,0),VLEN*sizeof (double),
          left_data_available[buffer_id], 1 + i, queue_id, GASPI_BLOCK));
    wait_or_die (segment_id, right_data_available[buffer_id], 1 + i);
    wait_or_die (segment_id, left_data_available[buffer_id], 1 + i);
}
#pragma omp barrier
data_compute ( NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
#pragma omp barrier
buffer_id = 1 - buffer_id;
```



# The GASPI Ring Exchange

- GASPI – left\_right\_double\_buffer\_multiple.c

```
if (tid == 0) {
    wait_for_queue_max_half (&queue_id);
    SUCCESS_OR_DIE ( gaspi_write_notify
        (segment_id, array_OFFSET_left (buffer_id, slice_id, 0), left,
         segment_id,array_OFFSET_left(buffer_id, right_halo,0),VLEN*sizeof(double),
         right_data_available[buffer_id], 1 + i, queue_id, GASPI_BLOCK));
    wait_or_die (segment_id, left_data_available[buffer_id], 1 + i);
    data_compute ( NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
} else if (tid < NTHREADS - 1) {
    data_compute ( NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
} else {
    wait_for_queue_max_half (&queue_id);
    SUCCESS_OR_DIE ( gaspi_write_notify
        ( segment_id, array_OFFSET_right (buffer_id, slice_id, 0), right,
         segment_id,array_OFFSET_right(buffer_id,left_halo, 0),VLEN*sizeof(double),
         left_data_available[buffer_id], 1 + i, queue_id, GASPI_BLOCK));
    wait_or_die (segment_id, right_data_available[buffer_id], 1 + i);
    data_compute ( NTHREADS, array, 1 - buffer_id, buffer_id, slice_id);
}
#pragma omp barrier
buffer_id = 1 - buffer_id;
```



Global  
Address Space  
Programming Interface  
**GASPI**

# Hands-On





# The GASPI Ring Exchange

- GASPI – `left_right_double_buffer_multiple.c`
  - One message instead of three (MPI Rendezvous)
  - No waiting for late MPI\_Recv
  - No waiting for acknowledge for MPI\_Isend
  - Overlap of communication with computation





# The GASPI Ring Exchange

- GASPI – Dataflow - left\_right\_dataflow\_halo.c

```
#pragma omp parallel default (none) firstprivate (buffer_id, queue_id) \
shared (array, data_available, ssl, stderr)
{
    slice* sl;
    while (sl = get_slice_and_lock (ssl, NTHREADS, num))
    {
        handle_slice(sl, array, data_available, segment_id, queue_id,
            NWAY, NTHREADS, num);
        sl->stage = sl->stage + 1;
        omp_unset_lock (&sl->lock);
    }
}
```

```
typedef struct slice_t
{
    omp_lock_t lock;
    volatile int stage;
    int index;
    enum halo_types halo_type;
    struct slice_t *left;
    struct slice_t *next;
} slice;
```



# The GASPI Ring Exchange

- GASPI – Dataflow - slice.c

```
void handle_slice ( ...)  
    if (sl->halo_type == LEFT){  
        if (sl->stage > sl->next->stage) {return;}  
        if (! test_or_die (segment_id, left_data_available[old_buffer_id], 1))  
            { return; }  
    } else if (sl->halo_type == RIGHT) {  
        if (sl->stage > sl->left->stage) { return; }  
        if (! test_or_die (segment_id, right_data_available[old_buffer_id], 1))  
            { return; }  
    } else if (sl->halo_type == NONE) {  
        if (sl->stage > sl->left->stage || sl->stage > sl->next->stage) {return;}  
    }  
data_compute (NTHREADS, array, new_buffer_id, old_buffer_id, sl->index);  
    if (sl->halo_type == LEFT) {  
        SUCCESS_OR_DIE ( gaspi_write_notify ...)  
    } else if (sl->halo_type == RIGHT)  
        SUCCESS_OR_DIE ( gaspi_write_notify ...)  
    }  
}
```



Global  
Address Space  
Programming Interface  
**GASPI**


# Hands-On





# GASPI – Dataflow

- Locally and globally asynchronous dataflow.





Global  
Address Space  
Programming Interface  
**GASPI**

# Collectives



# Collective Operations (I)

- Collectivity with respect to a definable subset of ranks (groups)
  - Each GASPI process can participate in more than one group
  - Defining a group is a three step procedure
    - `gaspi_group_create`
    - `gaspi_group_add`
    - `gaspi_group_commit`
  - `GASPI_GROUP_ALL` is a predefined group containing all processes



## Collective Operations (II)

- All gaspi processes forming a given group have to invoke the operation
- In case of a timeout (GASPI\_TIMEOUT), the operation is continued in the next call of the procedure
- A collective operation may involve several procedure calls until completion
- Completion is indicated by return value GASPI\_SUCCESS



# Collective Operations (III)

- Collective operations are exclusive per group
  - Only one collective operation of a given type on a given group at a given time
  - Otherwise: undefined behaviour
- Example
  - Two allreduce operations for one group can not run at the same time
  - An allreduce operation and a barrier are allowed to run at the same time





# Collective Functions

- Built in:
  - gaspi\_barrier
  - gaspi\_allreduce
    - GASPI\_OP\_MIN, GASPI\_OP\_MAX, GASPI\_OP\_SUM
    - GASPI\_TYPE\_INT, GASPI\_TYPE\_UINT,  
GASPI\_TYPE\_LONG, GASPI\_TYPE\_ULONG,  
GASPI\_TYPE\_FLOAT, GASPI\_TYPE\_DOUBLE
- User defined
  - gaspi\_allreduce user



# GASPI Collective Function

- **gaspi\_barrier**

```
gaspi_return_t  
gaspi_barrier ( gaspi_group_t group  
                , gaspi_timeout_t timeout )
```

- **gaspi\_allreduce**

```
gaspi_return_t  
gaspi_allreduce ( gaspi_const_pointer_t buffer_send  
                  , gaspi_pointer_t buffer_receive  
                  , gaspi_number_t num  
                  , gaspi_operation_t operation  
                  , gaspi_datatype_t datatype  
                  , gaspi_group_t group  
                  , gaspi_timeout_t timeout )
```



Global  
Address Space  
Programming Interface  
**GASPI**

# Passive communication



# Passive Communication Functions (I)

- 2 sided semantics send/recv
  - gaspi\_passive\_send

```
gaspi_return_t  
gaspi_passive_send ( gaspi_segment_id_t segment_id_local  
                    , gaspi_offset_t offset_local  
                    , gaspi_rank_t rank  
                    , gaspi_size_t size  
                    , gaspi_timeout_t timeout )
```

- time based blocking



# Passive Communication Functions (II)

## — Gaspi\_passive receive

```
gaspi_return_t  
gaspi_passive_receive ( gaspi_segment_id_t segment_id_local  
                        , gaspi_offset_t offset_local  
                        , gaspi_rank_t const *rank  
                        , gaspi_size_t size  
                        , gaspi_timeout_t timeout )
```

- Time based blocking
- Sends calling thread to sleep
- Wakes up calling thread in case of incoming message or given timeout has been reached



# Passive Communication Functions (III)

- Higher latency than one-sided comm.
  - Use cases:
    - Parameter exchange
    - management tasks
    - „Passive“ Active Messages (see advanced tutorial code)
      - GASPI Swiss Army Knife.



Global  
Address Space  
Programming Interface  
**GASPI**  
GASPI is a global address space programming interface for high performance computing.

# Fault Tolerance



# Features

- Implementation of fault tolerance is up to the application
- But: well defined and requestable state guaranteed at any time by
  - Timeout mechanism
    - Potentially blocking routines equipped with timeout
  - Error vector
    - contains health state of communication partners
  - Dynamic node set
    - substitution of failed processes





Global  
Address Space  
Programming Interface  
**GASPI**

# Questions?

Thank you for your attention

[www.gaspi.de](http://www.gaspi.de)

[www.gpi-site.com](http://www.gpi-site.com)